

# Astro\_FAQ · cmsa/Astronomer Wiki

[github.ford.com/cmsa/Astronomer/wiki/Astro\\_FAQ](https://github.com/ford/cmsa/Astronomer/wiki/Astro_FAQ)

cmsa

## 1. Even after removing DAGs, we could able to see it in the UI.

When you remove a DAG from your deployment, it should no longer be active. However, it might still appear in the Airflow UI for a couple of reasons:

1. Airflow scans for DAGs periodically: Airflow scans for new DAGs periodically (every minute by default). If you removed the DAG file very recently, it might still appear in the UI until the next scan.
2. DAG runs history: Even after a DAG is removed, its historical runs will still be visible in the Airflow UI. This is a feature of Airflow, not a bug. It allows you to see the history of your workflows, even after they've been removed or updated.
3. If you want to completely remove the DAG from the UI, you'll need to delete its records from the Airflow metadata database. Please be careful while doing this, as it involves directly manipulating the Airflow database and can lead to data loss if not done correctly. (Contact Astronomer support via <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/Astronomer-contact/>)
4. If the DAG is still appearing in the UI despite waiting for the next scan, it might be due to a caching issue, and you might need to clear the cache. If you're still having issues.

## 2. Is it possible to change dagbag\_import\_timeout from default value to 160 ?

Yes, by following this steps you can change the dagbag\_import\_timeout from default value to 160

1. Go to the Astronomer UI and navigate to the deployment you want to add the environment variable to.
2. Click on the Variables tab.
3. Click on the Add Variable button.
4. Enter the key-value pair of the environment variable you want to add. In this case add the environment variable `AIRFLOW__CORE__DAGBAG_IMPORT_TIMEOUT` and its value.
5. Click on the Add button.
6. After adding the environment variable, you'll need to redeploy your deployment for the changes to take effect.
7. For more information on environment values visit: <https://airflow.apache.org/docs/apache-airflow/stable/configurations-ref.html#dagbag-import-timeout>

3. Anyone here have worked on Identifying long running DAGs and send a notification via webex. ?

1. Please visit this wiki -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/airflow-alerting/> for setting up notification.

4. Can anyone assist me in understanding why, when I deploy my DAG changes through this Astro workspace, although the Tekton pipeline executes as expected and the last updated timestamp refreshes, the corresponding Docker tag does not update as it should? Additionally, why am I unable to locate the deployed DAG in Airflow?

1. Confirm if the Astro build is successful first and check whether the tags value is changed in Astro UI.
2. If build is successful and tag value is not updated in UI, please contact Astro support team using the below link -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/Astronomer-contact/>

Unable to locate the deployed DAGs in Airflow might be because of the following reasons.

1. DAG Parsing Issues: If there are syntax errors in your DAG file, Airflow might not be able to parse it correctly. Check your Airflow logs for any error messages related to DAG parsing.
2. Incorrect File Location: Make sure your DAG files are in the directory specified by the  `dags_folder`  configuration in your  `airflow.cfg`  file.
3. Folder Naming: If you provide the wrong names to the folder, Airflow won't be able to display it.

5. Anyone here have worked on setting up DAG run based on Event based. ? for example a pubsub notification to trigger cloud function --> cloud function will trigger an Airflow DAG or something similar using cloud run.

For event based trigger of Dag runs we can follow below methods:

1. TriggerDagRunOperator: Used when a system-event trigger comes from another DAG within the same Airflow environment.
2. Sensors: Sensors are useful when you need to coordinate the execution of tasks based on external events or dependencies.
3. Deferrable Operators: An option available to use when sensors, explained above, are ideal but the time of the system event is unknown. Deferrable operators are put in place so you don't have to leave a long-running sensor up all day, or forever, which would increase compute costs.
4. Airflow API: Used when the trigger event is truly random. It's the most reliable and low-cost method of monitoring system events in third-party applications outside of Airflow.

6. I am experiencing a problem with one of my tasks. When I attempt to view the logs, they are taking more than 15 minutes to load. Moreover, once they do load, I only see a small portion of the logs as the rest are truncated.

This issue could be due to a number of factors:

1. If your tasks are writing a large amount of logs, it could be causing delays in fetching and displaying them in the UI. Consider reducing the verbosity of your logging if possible.
2. If you are using remote logging (GCP Bucket, PubSub), there could be network latency or issues with the remote logging service. Ensure your remote logging service is properly configured and functioning correctly.
3. If your Airflow webserver is under heavy load, it could be slowing down the loading of logs. Consider scaling up your webserver or reducing its load if possible. As for the truncated logs, Airflow by default only fetches a limited amount of log data to display in the UI.

7. Could anyone please assist me in resolving an error I'm encountering with a Dataproc task? The error message I'm receiving is: "ERROR - Failed to execute job for task dataproc\_preprocessing (409 IP space of 'projects//regions/' exhausted)." How can I address this issue of IP space exhaustion in my GCP project?

1. If you face any IP exhaustion issues in the dataproc service, please verify whether the VLANs provided in the dataproc repository is correct.
2. For further support please contact the dataproc service team. Link-> <https://docs.gcp.ford.com/docs/services/ai---ml---analytics/dataproc/>

8. Is there any document to create Astonomer Local Environment(Astro CLI) with Dbt Core.

1. You can follow the below tutorial to set Astronomer Locally -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/astro-localdev/>
2. You can find the provider packages and related operators in [registry.astronomer.io](https://registry.astronomer.io)

9. Hi Team, we are trying to create a DAG to copy a file from one GCS bucket to different folder in same bucket. Getting error as "Broken Dag: Syntax error" while deploying the script. Could someone please help on this ?

1. Please use the following documentation and use the suggested airflow provider operator GCSFileTransformOperator to copy a file from one folder to another in the same Google Cloud Storage (GCS) bucket. Source Link -> [https://airflow.apache.org/docs/apache-airflow-providers-google/stable/\\_api/airflow/providers/google/cloud/transfers/gcs\\_to\\_gcs/index.html](https://airflow.apache.org/docs/apache-airflow-providers-google/stable/_api/airflow/providers/google/cloud/transfers/gcs_to_gcs/index.html)
2. Syntax errors can occur for a variety of reasons, such as:
  1. Misspelled keywords or identifiers
  2. Missing or misplaced punctuation (e.g., commas, semicolons, braces)
  3. Incorrect use of operators
  4. Improper structure or indentation (especially in languages like Python that rely on indentation)
  5. Incorrect function or method calls
  6. Mismatched parentheses, brackets, or braces
3. You can find the complete list of airflow component and operators in [registry.astronomer.io](https://registry.astronomer.io)

10. How to create an image pull secret for Ford Registry and add it to Astro Tekton pipeline? //

Hey Guys ... I am having this issue with image registry: failed to fetch anonymous token.

"ERROR: failed to solve: registry.ford.com/astro/astro-runtime:6.0.3: failed to authorize: failed to fetch anonymous token: unexpected status from GET request to <https://registry.ford.com/v2/auth?scope=repository%3Aastro%2Fastro-runtime%3Apull&service=registry.ford.com>: 401 Unauthorized Error: command 'docker build -t milo\_09da08/airflow:latest failed: failed to execute cmd: exit status 1 "

can anyone help?

Check whether the image pull secret is created and the correct secret is used in Tekton. Steps to create an image pull secret and add in Tekton is as follows:

1. Login to the registry.ford.com, create a robot account in your organisation.
2. If you are not part of any organization, then you need to create an organization.
3. Once you are in your organization registry, create robo account.
4. Once robot account is created, click on the account.
5. Click on kubernetes secret
6. Click on view secret.yaml
7. Copy the yaml.
8. Below is the example for the pull-secret

```
apiVersion: v1 kind: Secret metadata: name: image-pull data: .dockerconfigjson:
XXXXXXXX type: kubernetes.io/dockerconfigjson
```

9. For creating secrets in the tekton namespace click on workloads
10. Click on secrets
11. Click on create
12. Click on copy from yaml and paste the yaml
13. Click on Create

11. Hi Team,

i'm getting following error at ftp\_sensor step , please help me resolve this.

[\*\*\* Log DUAL\_AUTH\_CDS\_SFTGDIA1\_FORD\_CDS\_AOD\_PACKLIST-ftp-sensor-manual\_\_2024-01-08T18:35:34.033010+00:00--1-1 not found in Elasticsearch. If your task started recently, please wait a moment and reload this page. Otherwise, the logs for this task instance may have been removed.]

1. Try retriggering the DAGs after ensuring all the environmental variables, connections and dependencies are available to them along with the requisite resources by analysing the metrics tab.
2. If the issue persists, please contact Astronomer support by following the below link -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/Astronomer-contact/> Note for Astro-support team: Verify all the Airflow Components and StatefulSets are Up and running for Kubernetes/Celery Executors.

12. Hi team,

Could you please help me to resolve the below issue occurred during the dag run.

I ran a dag which should execute series of SQL files but failed with below log,

\*\*\* Log IDM\_01\_ACCT\_PROCESSstg-na\_scaid\_guid\_lkup\_tbl\_crt-manual\_\_2024-01-08T05 58.048743+00:00--1-1 not found in Elasticsearch. If your task started recently, please wait a moment and reload this page. Otherwise, the logs for this task instance may have been removed.

1. Try retriggering the DAGs after ensuring all the environmental variables, connections and dependencies are available to them along with the requisite resources by analysing the metrics tab.
2. If the issue persists, please contact Astronomer support by following the below link -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/Astronomer-contact/> Note for Astro-support team: Verify all the Airflow Components and StatefulSets are Up and running for Kubernetes/Celery Executors.

13. Dags are failing or not running without any proper logs available in the ElasticSearch.

1. Try retriggering the DAGs after ensuring all the environmental variables, connections and dependencies are available to them along with the requisite resources by analysing the metrics tab.
2. If the issue persists, please contact Astronomer support by following the below link -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/Astronomer-contact/> Note for Astro-support team: Verify all the Airflow Components and StatefulSets are Up and running for Kubernetes/Celery Executors.

14. Unable to see workspace / Customer Astro workspace not visible or listed in FCP portal.

1. Log into the environment for which you are facing the issue in Astro platform
2. Open the workspace view and click on the "Teams" tab
3. Review the workspace owner group, and make sure the role is "workspace admin"
4. Make sure you are part of the FIM group.
5. However if you don't see even a workspace, please raise an it connect ticket w/ the astro team including the above reviewed details.

15. Hi Team, GWTS dev pipeline failing in dataproc step due to below error, Kindly look into this. IP address exhaustion error. google.api\_core.exceptions.Aborted: 409 IP space of 'projects//regions/' is exhausted. 10: IP space of 'projects//regions/' is exhausted.

1. If you face any IP exhaustion issues in the GCP service, please verify whether the VLANs provided in the GCP Service repository is correct.
2. For further support please contact the GCP service team. Link-> <https://docs.gcp.ford.com/docs/services/>

## 16. What are the reasons for dags not able to show in the airflow?

1. Folder Naming Standard: If you provide the wrong names to the folder, Airflow won't be able to display it. Please follow the same format and namings provided in the document.
2. DAG values: Dag id values should be unique.
3. DAG Parsing Issues: If there are syntax errors in your DAG file, Airflow might not be able to parse it correctly. Check your Airflow logs for any error messages related to DAG parsing.
4. Incorrect File Location: Make sure your DAG files are in the directory specified by the `dags_folder` configuration in your `airflow.cfg` file.
5. Scheduler is Not Running: The Airflow scheduler needs to be running for DAGs to show up in the UI.
6. DAG Bag Loading Issues: If there are too many DAGs or if DAG files are too large, it might take some time for Airflow to load all of them. You can check the `dagbag_size` and `dagbag_import_timeout` configurations in your `airflow.cfg` file.
7. Incorrectly Assigned DAG: Every Operator/Task must be assigned to a DAG to run. If you forget to assign an Operator to a DAG, Airflow won't be able to display it.

## 17. The astronomer deployment Current Tag is not changing after pushing new changes. What might be the cause?

The tag shown in the Airflow UI is a unique identifier generated by a CI tool or a timestamp generated by the Astro CLI on `astro deploy`. If it is not changing after pushing new changes, it may indicate that the changes were not successfully deployed.

Here are a few things you can do to troubleshoot this:

1. Verify that your astro deployment completed successfully without any error logs in the tekton pipeline. If still facing the issue, contact the astronomer team

## 18. Dag is failing at random and not get proper logs for the failed dags.

1. verify the deployment resources utilization under the "metric's" section in the UI to confirm airflow resources components are not over utilized
2. Increase the resources if it's overused.
3. implement the retry logics in the task to avoid DAG failures while execution.

## 19. Deployment is in unknown state and airflow is showing rocket page.

1. If deployment is in unknown state: Triggering the pipeline can bring the deployment back to online.
2. If deployment is in healthy state and airflow page is showing the rocket/ error page follow the below steps:
  - Check if you have provisioned adequate resources (CPU, Scheduler resources etc) by reviewing the metrics tab in the Astronomer UI of the deployment.
  - If any recent code changes are made, revert back the changes and check the status of the deployment.
  - If you have set this variable "AIRFLOW\_\_WEBSERVER\_\_WORKERS" remove the variable from astronomer deployment UI.
  - Check if you have provided the required environmental variables and connections to your deployment.
  - If you have recently provisioned the deployment, kindly wait for at least 15 minutes for the Airflow UI to show up.
  - If you are still facing issue, kindly raise a support ticket with Astronomer Support team to check the backend.

## 20. what are the reasons for a deployment to be in unhealthy state?

A deployment can be in an unhealthy state due to a variety of reasons:

1. Resource constraints: If the resources allocated to your Airflow deployment (CPU, memory, storage) are not sufficient, it can cause your tasks to fail and put your deployment in an unhealthy state.
2. Task scheduling latency: If the scheduler does not have enough resources to parse DAGs and schedule tasks, it can cause high latency and result in an unhealthy deployment.
3. DAGs in queued state: If the number of tasks being scheduled is beyond the capacity of your Airflow infrastructure, DAGs may remain in a queued state and not run.
4. Database issues: Problems with your database, such as memory issues, connectivity problems, or configuration errors, can cause your deployment to be unhealthy.

## 21. Astronomer deployment in progress since long time.



1. Check the resources usage in “CPU usage” under “Historical” which is present in astronomer deployment “metric”.
2. Increase the resources if its overused.
3. If any recent changes made in code/dockerfile version change/ variables and other configuration items and confirm if the build/pipeline is successful.
4. In the deployment metric tab section, please check whether the core airflow components (webserver, triggerer, pgbouncer etc) are in healthy state. (green)
5. You can check the live logs of your deployment and its components in the UI for debugging the issue.
6. If not contact the astronomer team

22. One of our deployments is in a constant deployment state even though we havent deployed recently. could someone help resolve that?

1. Please wait for 15 minutes for all the pods to be up and deployment to be in healthy state in the Astro UI.
2. In the deployment metric tab section, please check whether the core airflow components (webserver, triggerer, pgbouncer etc) are in healthy state. (green)
3. You can check the live logs of your deployment and its components in the UI for debugging the issue.
4. If you are still facing some issue, please raise a ticket to Astro support team for further investigation and support.

23. Hello Team, I tried deploying new changes to astro with a new version but looks like the tag is not updated in the deployment. Could you please look into this and help me with the resolution?

1. Confirm if the Astro build is successful first and check whether the tags value is changed in Astro UI.
2. If build is successful and tag value is not updated in UI, please contact Astro support team using the below link ->  
<https://pages.github.ford.com/gcam/astro-docs/astro/Astronomer-contact/>

24. Hi Team, is there any issue with Airflow? Our dag failed in dev with error msg - Exception: Error in BigQuery Procedure: Error encountered during execution. Retrying may solve the problem. Yesterday, we faced similar issue in prod as well. We didn't find any error in log as well

1. Try retriggering the DAGs after ensuring all the environmental variables, connections and dependencies are available to them along with the requisite resources by analysing the metrics tab.
2. If the issue persists, please contact Astronomer support by following the below link -> <https://pages.github.ford.com/gcam/astro-docs/astro/Astronomer-contact/>

25. Hi Team , I m getting this error "Error: Error waiting to create Service: resource is in failed state "Ready:False", message: Revision 'df-dataingest-cloudrun-astro-trigger-dag-00001-7p9' is not ready and cannot serve traffic.

spec.template.spec.containers[0].env[5].value\_from.secret\_key\_ref.name: Secret projects/712389963333/secrets/astro-token-qa/versions/2 was not found " while configuring the Cloud Run. // Unable to fetch, access secrets or tokens.

1. Please verify that the secrets are correctly configured in the GCP Secret Manager and in the Astronomer Deployment GCP Connection Configurations with the appropriate values.
2. Additionally, ensure that the Astronomer Deployment Service account has the necessary permissions for GCP services.

26. Hi team, We are trying to push the code in the dev and the tekton pipeline is failing with the below error: container "step-get-astro-config" in pod "astro-ci-bp-dfd010af5245ca22db52d713df79b10c9df0c058418cd86-pod" is waiting to start: CreateContainerConfigError

1. Check whether the secrets are created and configured for workspace, image pull, image push.
2. Also ensure the secrets are patched with the pipeline.

27. How can we encrypt password and other data In astronomer variable?

1. You can use Airflow Secrets Backend to store sensitive connections and variables in a secure manner using GCP Secret Manager.
2. Refer this doc on how to set up secrets backend.
  - a. [https://pages.github.ford.com/gcam/astronomer-docs/astronomer/astro\\_external\\_secrets\\_backend/](https://pages.github.ford.com/gcam/astronomer-docs/astronomer/astro_external_secrets_backend/)
  - b. <https://docs.astronomer.io/astro/secrets-backend>
3. Astronomer uses a Kubernetes Secret to store environment variables added via the Astronomer UI. These variables are encrypted at rest and are automatically mounted to your Airflow pods.

28. Hi Support Team, Latest Changes to Deployment "Dev common" not reflecting. can you please help? // Chnages to deployment not reflected.

1. Check if the pipeline is completed, Astro build is success and Tag is updated.
2. Wait for the deployment to ingest the DAGs / check for DAG import errors
3. Check if the latest code is reflected in the Code tab of the DAG
4. If you are unable the view the DAG, check if it has a unique name. DAGs should have a unique name across deployment.
5. Try the above steps and if you are still facing issue, please raise a ticket with the Astro support team to help you resolve it.

29. Hi Team, Our deployment CI pipeline got broke. Any one in this forum can help us get this up. // Pipeline broke or stuck or not working.

1. If you are using the standard Astronomer CI pipeline, please raise a ticket with Astro Support team for further investigation and support.
2. If you are using a custom pipeline, please contact your respective internal teams to resolve the issue.

30. How would i go about getting the logs from an astronomer job that has failed but no longer outputs any logs?

1. Try retriggering the DAGs after ensuring all the environmental variables, connections and dependencies are available to them along with the requisite resources by analysing the metrics tab.
2. Try refereshing the page
3. If the issue persists, please contact Astronomer support by following the below link -> <https://pages.github.ford.com/gcam/astro-docs/astro/Astronomer-contact/> Note for Astro-support team: Verify all the Airflow Components and StatefulSets are Up and running for Kubernetes/Celery Executors.

31. Hi Team, We are facing the Tekton issue. Fossa and CheckMarx failed, for Astro. Even the Pipeline configuration is marked as "enabled": false. Kindly guide us to resolve this issue.

1. If you are using the standard Astronomer CI pipeline, please raise a ticket with Astro Support team for further investigation and support.
2. If you are using a custom pipeline, please contact your respective internal teams to resolve the issue.
3. Fossa and CheckMarx are custom pipelines ,Astro team is not handling these pipelines . Kindly check with the respective team.

32. Hi Team I'm facing issues in importing gcs\_bq.py DAG to my astro workspace.

This is repo linked to my Astro workspace and I'm using this [registry.ford.com/gdia-datafactory/astro-dfingesttools:5.0.11-10.03-dev](https://registry.ford.com/gdia-datafactory/astro-dfingesttools:5.0.11-10.03-dev) as astro base image. Can someone please help if I'm missing something? TIA

1. If you are using customised images, please check with the image owners and respective teams for import/code related issues
2. If you are using astro official images, please check if you have included all the requirements in the requirements.txt and whether you are using the correct import operators and provider package versions.
3. You can visit [registry.astronomer.io](https://registry.astronomer.io) for operator documentation and examples.
4. If you require additional support when using astronomer official images, please contact astronomer support team via a ticket.

33. hello, I have a customer stating their Astronomer Workspaces and Deployments were not created when project m1-mlops-fedebom was provisioned through mach1ml. can I get some assistance/support on this please?

Please refer Ford Astro documentation on how to provision a workspace and a deployment through FCP and ensure if you have followed them correctly. Link-> <https://pages.github.ford.com/gcam/astro-docs/astro/Astronomer-FCP/>

34. Hi Team , while migrating the astro code in QA (service360-qa) , getting following error "container "step-fossa" in pod "astro-ci-gdia-service360-astro-service360-gcp-xg6911-fossa-pod" is waiting to start: PodInitializing" . Need your support in resolving issue.

1. If you are using the standard Astronomer CI pipeline, please raise a ticket with Astro Support team for further investigation and support.
2. If you are using a custom pipeline, please contact your respective internal teams to resolve the issue.
3. check\_Fossa is a custom pipeline task ,Astro team is not handling these pipelines . Kindly check with the respective team.
4. Verify tekton resource has any capacity constraints in namespace.

35. Hey, is there any way to retrieve or regenerate a workspace SA key? I dont think we got the email from FCP with the key attached. Thanks!

Please contact Astro Support Team via a Ticket for help on this issue.

36. Hi Team any updates on DEV version of Astronomer RunTime Version 5.4.0 as both QA and PROD we have 5.0.11. Thanks !

The latest version of Astronomer Runtime is 9.7.0.

37. Hi Team, I tried to upgrade from runtime 8.1.0 to 9.7.0, astro-build did break due to some other package dependency. Then, I am tryin to revert to 8.1.0, I changed the version on FCP/astro --> update deployment. Then, pushed the docker change to make it 8.1.0. Now, when I push PR, changes are not being deployed and I see this upgrade attempt still on the deployment page.

Can some one reset version to 8.1.0 in the backend please?

1. Currently rollback or downgrade of Astronomer runtime from newer to older version is not supported.
2. Please contact Asrtronomer support team with a ticket for further support on this issue.

38. Hello Team - Need assistance to move the astro pipeline to active - REQ000006910723

1. Please engage with the Astro support engineer assigned to this ticket.
2. Our team will look into the ticket/request and get back to you.

39. Hi, For CI configuration/ DAG code, where do we need to create github repository? is it on on-prem (github.ford.com) or cloud (github.com). Please advise

1. User should create Astro CI repository in github.ford.com (on-prem).

40. Hi Team, I am facing issue in Astro notify-start in Astro Tekton pipeline. But same github-https secret is working for my Terraform Tekton pipeline <https://console-openshift-console.apps.pd303.caas.gcp.ford.com/pipelines/ns/dfdm-84-gfabc/pipeline-runs>

1. Verify the Git Hub HTTP web hook secret and tekton pipeline github-https secret are same and retrigger the pipeline.
2. Verify if the secrets were created as per the specified format and were imported into tekton using import from YAML option as mentioned in the pipeline setup documentation.

41. hello, we recently created a remote astronomer workspace (following these steps: [https://pages.github.ford.com/Mach1ML-2-0/MDK-Docs/Getting\\_Started/](https://pages.github.ford.com/Mach1ML-2-0/MDK-Docs/Getting_Started/))

The PR has been merged but we aren't able to login to this provided link to verify the workspace.

This is the error seen when logging in with the ford login:

[https://app.astrodev.ford.com/login?error=PUBLIC\\_SIGNUPS\\_DISABLED](https://app.astrodev.ford.com/login?error=PUBLIC_SIGNUPS_DISABLED)

1. Please follow this documentation to create the astronomer workspace and deployments using Astronomer tile in ford cloud portal. Link -> <https://pages.github.ford.com/gcam/astrodev-docs/astrodev/Astronomer-FCP/>

42. Hi Team, Getting the below error while running the gcs\_bq dag [2024-01-26, 09:10:36 UTC] {standard\_task\_runner.py:92} ERROR - Failed to execute job 7802 for task build\_stg\_config (local variable 'load\_stg\_config' referenced before assignment; 39)

1. Check if the environment variables is present in your airflow deployment.
2. To verify, go to airflow UI, click on Admin.
3. Click on Variables.
4. Click on import/create the environmental variable, referenced by your dag. Please follow Astro documentation on creating environment variables. Link-> <https://pages.github.ford.com/gcam/astrodev-docs/astrodev/astro-variables-connections/#astrodev-variables>

43. Hi Team, I created a deployment in dev last night via FCP/astronomer self-service, I don't see it under dev workspace, can some one look into it please?

1. Check if you have received the workspace created success confirmation message with workspace UI links after submitting all the required deployment.
2. Please login to the UI and verify if you can see the deployments.
3. Please verify if you are part of the FIM groups assigned to the workspace.
4. Please contact support if you are still facing issue after following the above steps.

44. I am having trouble logging into the astronomer registry from powershell, getting error that the "service unavailable".. I am able to login to harbor.hpc.ford.com and application registry, but for some reason astro registry isnt working

1. Verify if you are in the ford network and are part of your organization in the registry.
2. Verify if all the required secrets and credentials are available and configured for image pull/image push etc.

45. runtime updated in cloud portal , but it says Upgrade from 5.0.11 to 8.1.0 initiated... , please check

<https://app.astrodev.ford.com/w/cli25brgz27096114hpfy4ubswb/d/infrared-waxing-4665>

1. Please follow the documentation to upgrade Astro Run time. Link -> <https://pages.github.ford.com/gcam/astronomer-docs/astronomer/astro-runtime-customer-doc/Readme/#steps-to-be-followed-for-migration-to-runtime>
2. In addition to making update in FCP you need to make changes to the docker file and raise a PR to trigger the upgrade in Astro platform.

46. Hi, I am trying to deploy an Airflow DAG in astronomer and when we try to create a deployment workspace we get the error in screenshot. The person trying to create the deployment is the owner of the FIM group we created for this project. Is my understanding of "should be part of owner delgate FIM groups" right? Do you have any suggestions on how to resolve this?

1. The person trying to create the deployment must be part of the FIM group assigned to Astronomer workspace
2. The FIM group should start with astro & contain astrodev user.
3. If you have recently added the member to the FIM group, please wait upto 2 hours for the changes to propagate.
4. User has to login to the respective dev, prod, preprod Astronomer UI using browser to get synced with Astronomer.
5. Renaming the existing FIM group names (including case sensitive changes) to make them compliant with above rules, will not work. User has to create a new FIM group compliant with the above rules.

47. Hi, We tried increasing the scheduler from 1 to 2. However the scheduler memory is still showing in 100%. Once GCS\_BQ is triggered it is moving from queued to failed stage and not showing any logs.

Please advise on this asap. As its delaying our production.

1. Check the metrics tab under the deployment in Astro UI, and increase the resources.
2. Try increasing the memory, CPU for the scheduler and check utilization after 15 minutes.
3. Monitor the DAG resource consumption using the metrics tab and provision adequate resources.

48. Hi team. I'm new to using Astronomer and I'm currently in the process of creating a workspace. I have a question regarding the FIM group. Is it a requirement for the FIM group to start with "astro-"? I currently have an FIM group that starts with "gcp-". Will this FIM group still work with Astronomer?

1. Astronomer FIM group should start with astro- prefix only
2. Astronomer FIM group should contain astrodev user in it.

49. Hi all, Our astro deployment has been facing some issues which seem like other teams have experienced before. We are unable to view the logs for a failed task and the UI just displays - \*\*\* Log fpm\_training\_validation-task\_group\_training.file\_transfer-manual\_\_2024-02-05T14:54:03.506957+00:00--1-1 not found in Elasticsearch. If your task started recently, please wait a moment and reload this page. Otherwise, the logs for this task instance may have been removed.

We were also facing the scheduler/heartbeat error last week.

The deployment link is - <https://deployments.astrodev.ford.com/heliocentric-axis-8703/airflow/home>

1. Try retriggering the DAGs after ensuring all the environmental variables, connections and dependencies are available to them along with the requisite resources by analysing the metrics tab.
2. If the issue persists, please contact Astronomer support by following the below link -> <https://pages.github.ford.com/gcam/astrodev/docs/astrodev/Astronomer-contact/> Note for Astro-support team: Verify all the Airflow Components and StatefulSets are Up and running for Kubernetes/Celery Executors.

50. We are seeing that our prod dag is failing with the following exception: Task Instance Running as Zombie. Please let us know how can we resolve these kind of exceptions. It keeps happening once in a while. How to resolve the exceptions that says caused by "Task Instances running as Zombie"

Zombie tasks in Airflow are tasks that are supposed to be running but have suddenly died (e.g. their process was killed, or the machine died). Airflow periodically checks for these tasks, cleans them up, and either fails or retries the task depending on its settings. Here are some steps you can take to resolve issues with zombie tasks:

**Increase the Heartbeat Interval:** You can increase the `scheduler_zombie_task_threshold` configuration value in your `airflow.cfg` file. This is the time in seconds after which a task would be considered a zombie.

**Check the System Resources:** Make sure that your system has enough resources (CPU, Memory). If a system is running out of resources, it might kill processes leading to zombie tasks.

**Check the Logs:** The logs can provide valuable information about why a task was killed. You might find that the task was consuming too much memory or there was an unhandled exception in your task code.

**Retry Failed Tasks:** If tasks are failing, you can set retries and `retry_delay` parameters in your tasks to make Airflow retry failed tasks.

**Check your Task Code:** If a task code has issues, it might be causing the task to fail and become a zombie. Make sure your task code is correct and can handle exceptions properly.

**Check Worker Availability:** If you are using CeleryExecutor, make sure your workers are running and are not overloaded with tasks. If a worker dies while processing a task, the task might become a zombie. Remember, identifying the root cause of zombie tasks can be complex and might require a thorough investigation of your Airflow setup, task definitions, and system resources. Please raise a ticket with astro support if the above steps are not working.

51. Hi Team.. Getting 504 Gateway time-out when tried to open the airflow ui <https://deployments.astro.ford.com/pure-ion-0025/airflow/>. Is there any resolution for this?



1. If deployment is in unknown state: Triggering the pipeline can bring the deployment back to online.
2. If deployment is in healthy state and airflow page is showing the rocket/ error page follow the below steps:

- Check if you have provisioned adequate resources (CPU, Scheduler resources etc) by reviewing the metrics tab in the Astronomer UI of the deployment.
- If any recent code changes are made, revert back the changes and check the status of the deployment.
- If you have set this variable "AIRFLOW\_\_WEBSERVER\_\_WORKERS" remove the variable from astronomer deployment UI.
- Check if you have provided the required environmental variables and connections to your deployment.
- If you have recently provisioned the deployment, kindly wait for at least 15 minutes for the Airflow UI to show up.
- If you are still facing issue, kindly raise a support ticket with Astronomer Support team to check the backend.

52. Hi Team, our deployment also unhealthy and it shows : The scheduler does not appear to be running. Last heartbeat was received 14 hours ago.

1. If deployment is in unknown state: Triggering the pipeline can bring the deployment back to online.
2. If deployment is in healthy state and airflow page is showing the rocket/ error page follow the below steps:

- Check if you have provisioned adequate resources (CPU, Scheduler resources etc) by reviewing the metrics tab in the Astronomer UI of the deployment.
- If any recent code changes are made, revert back the changes and check the status of the deployment.
- If you have set this variable "AIRFLOW\_\_WEBSERVER\_\_WORKERS" remove the variable from astronomer deployment UI.
- Check if you have provided the required environmental variables and connections to your deployment.
- If you have recently provisioned the deployment, kindly wait for at least 15 minutes for the Airflow UI to show up.
- If you are still facing issue, kindly raise a support ticket with Astronomer Support team to check the backend.

53. We have a 503 error on our deployment, could anyone help look into this?

<https://deployments.astropp.ford.com/frozen-proton-1870/airflow/>

1. If deployment is in unknown state: Triggering the pipeline can bring the deployment back to online.
2. If deployment is in healthy state and airflow page is showing the rocket/ error page follow the below steps:

- Check if you have provisioned adequate resources (CPU, Scheduler resources etc) by reviewing the metrics tab in the Astronomer UI of the deployment.
- If any recent code changes are made, revert back the changes and check the status of the deployment.
- If you have set this variable "AIRFLOW\_\_WEBSERVER\_\_WORKERS" remove the variable from astronomer deployment UI.
- Check if you have provided the required environmental variables and connections to your deployment.
- If you have recently provisioned the deployment, kindly wait for at least 15 minutes for the Airflow UI to show up.
- If you are still facing issue, kindly raise a support ticket with Astronomer Support team to check the backend.

54. Hi Team, How do I install the following package apache-airflow-macros? I would like to use the following macro in a script/SQL template `airflow.macros.ds_add`

For information on installing packages, provider packages, using operators and documentation with example / sample code you can visit [registry.astronomer.io](https://registry.astronomer.io).

The `airflow.macros.ds_add` is a built-in macro in Apache Airflow and does not require any additional package installation. It's part of the core Airflow distribution. You can use it in your Airflow templates like this: `{{ macros.ds_add(ds, 1) }}`. This will add one day to the ds (execution date in YYYY-MM-DD format). For more information on using macros, refer to the <https://airflow.apache.org/docs/apache-airflow/stable/templates-ref.html#macros>